

▼ 1. Keras – MLPs on MNIST Assignment

▼ 1.1 Importing required Libraries

```
# if you keras is not using tensorflow as backend set "KERAS_BACKEND=tensorflow" u
import tensorflow as tf
from keras.utils import np_utils
from keras.datasets import mnist
import seaborn as sns
from keras.initializers import RandomNormal
from keras.layers.normalization import BatchNormalization
from keras.layers.normalization import BatchNormalization
from keras.layers import Dropout
```

➞ The default version of TensorFlow in Colab will soon switch to TensorFlow 2.x.
We recommend you [upgrade](#) now or ensure your notebook will continue to use TensorFlow 1.x via the %t
1.x magic: [more info](#).
Using TensorFlow backend.

▼ 1.2 Function to plot a dynamic plot

```
import matplotlib.pyplot as plt
%matplotlib inline
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()
```

▼ 1.3 High level overview of data set

```
# the data, shuffled and split between train and test sets
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

➞ Downloading data from <https://s3.amazonaws.com/img-datasets/mnist.npz>
11493376/11490434 [=====] - 2s 0us/step

```
print("Number of training examples :", X_train.shape[0], "and each image is of sha
```

```
print("Number of training examples :", X_test.shape[0], "and each image is of shap
```

```
↳ Number of training examples : 60000 and each image is of shape (28, 28)
   Number of training examples : 10000 and each image is of shape (28, 28)
```

```
# if you observe the input shape its 2 dimensional vector
# for each image we have a (28*28) vector
# we will convert the (28*28) vector into single dimensional vector of 1 * 784
```

```
X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

```
# after converting the input images from 3d to 2d vectors
```

```
print("Number of training examples :", X_train.shape[0], "and each image is of sha
print("Number of training examples :", X_test.shape[0], "and each image is of shap
```

```
↳ Number of training examples : 60000 and each image is of shape (784)
   Number of training examples : 10000 and each image is of shape (784)
```

```
# An example data point
print(X_train[0])
```

```
↳
```

```
[
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  3  18  18  18 126 136 175 26 166 255
247 127  0  0  0  0  0  0  0  0  0  0  0  0  0  30 36 94 154
170 253 253 253 253 253 225 172 253 242 195 64  0  0  0  0  0  0
  0  0  0  0  0  49 238 253 253 253 253 253 253 253 251 93 82
82 56 39  0  0  0  0  0  0  0  0  0  0  0  0  18 219 253
253 253 253 253 198 182 247 241  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0 80 156 107 253 253 205 11  0 43 154
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0 14  1 154 253 90  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0 139 253 190  2  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0 11 190 253 70  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0 35 241
225 160 108  1  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0 81 240 253 253 119 25  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0 45 186 253 253 150 27  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0 16 93 252 253 187
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0 249 253 249 64  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0 46 130 183 253
253 207  2  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0 39 148 229 253 253 253 250 182  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0 24 114 221 253 253 253
253 201 78  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0 23 66 213 253 253 253 253 198 81  2  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0 18 171 219 253 253 253 253 195
80  9  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
55 172 226 253 253 253 253 244 133 11  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0 136 253 253 253 212 135 132 16
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
]
```

▼ 1.4 Normalizing the train and test sets

```
# if we observe the above matrix each cell is having a value between 0-255
# before we move to apply machine learning algorithms lets try to normalize the da
# X => (X - Xmin)/(Xmax-Xmin) = X/255
```

```
X_train = X_train/255
X_test = X_test/255
```

```
# example data point after normlizing
print(X_train[0])
```

File Edit Insert View Help



1

▼ 1.5 One-Hot Encoding the class label

```
# here we are having a class number for each image
print("Class label of first image :", y_train[0])

# lets convert this into a 10 dimensional vector
# ex: consider an image is 5 convert it into 5 => [0, 0, 0, 0, 0, 1, 0, 0, 0, 0]
# this conversion needed for MLPs

Y_train = np_utils.to_categorical(y_train, 10)
Y_test = np_utils.to_categorical(y_test, 10)

print("After converting the output into a vector : ",Y_train[0])

☞ Class label of first image : 5
   After converting the output into a vector : [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]
```

▼ 1.6 A simple 2 layer model with Softmax classifier

```
# https://keras.io/getting-started/sequential-model-guide/

# The Sequential model is a linear stack of layers.
# you can create a Sequential model by passing a list of layer instances to the co

# model = Sequential([
#     Dense(32, input_shape=(784,)),
#     Activation('relu'),
#     Dense(10),
#     Activation('softmax'),
# ])

# You can also simply add layers via the .add() method:

# model = Sequential()
# model.add(Dense(32, input_dim=784))
# model.add(Activation('relu'))

####

# https://keras.io/layers/core/
```

```

# keras.layers.Dense(units, activation=None, use_bias=True, kernel_initializer='gl
# bias_initializer='zeros', kernel_regularizer=None, bias_regularizer=None, activi
# kernel_constraint=None, bias_constraint=None)

# Dense implements the operation: output = activation(dot(input, kernel) + bias) w
# activation is the element-wise activation function passed as the activation argu
# kernel is a weights matrix created by the layer, and
# bias is a bias vector created by the layer (only applicable if use_bias is True)

# output = activation(dot(input, kernel) + bias) => y = activation(WT. X + b)

####

# https://keras.io/activations/

# Activations can either be used through an Activation layer, or through the activ

# from keras.layers import Activation, Dense

# model.add(Dense(64))
# model.add(Activation('tanh'))

# This is equivalent to:
# model.add(Dense(64, activation='tanh'))

# there are many activation functions ar available ex: tanh, relu, softmax

from tensorflow.python.keras.layers import Dense
from tensorflow.python.keras.layers import BatchNormalization
from tensorflow.python.keras import Sequential

# some model parameters

output_dim = 10
input_dim = X_train.shape[1]

batch_size = 256
nb_epoch = 20

# start building a model
model = Sequential()

# The model needs to know what input shape it should expect.
# For this reason, the first layer in a Sequential model
# (and only the first, because following layers can do automatic shape inference)
# needs to receive information about its input shape.
# you can use input_shape and input_dim to pass the shape of input

# output_dim represent the number of nodes need in that layer
# here we have 10 nodes

model.add(Dense(output_dim, input_dim=input_dim, activation='softmax'))

```



```

↳ WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow_core/
Instructions for updating:
If using Keras pass *_constraint arguments to layers.

```

```

# Before training a model, you need to configure the learning process, which is do

```

```

# It receives three arguments:

```

```

# An optimizer. This could be the string identifier of an existing optimizer , htt

```

```

# A loss function. This is the objective that the model will try to minimize., htt

```

```

# A list of metrics. For any classification problem you will want to set this to m

```

```

# Note: when using the categorical_crossentropy loss, your targets should be in ca

```

```

# (e.g. if you have 10 classes, the target for each sample should be a 10-dimensio

```

```

# for a 1 at the index corresponding to the class of the sample).

```

```

# that is why we converted out labels into vectors

```

```

model.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy

```

```

# Keras models are trained on Numpy arrays of input data and labels.

```

```

# For training a model, you will typically use the fit function

```

```

# fit(self, x=None, y=None, batch_size=None, epochs=1, verbose=1, callbacks=None,

```

```

# validation_data=None, shuffle=True, class_weight=None, sample_weight=None, initi

```

```

# validation_steps=None)

```

```

# fit() function Trains the model for a fixed number of epochs (iterations on a da

```

```

# it returns A History object. Its History.history attribute is a record of traini

```

```

# metrics values at successive epochs, as well as validation loss values and valid

```

```

# https://github.com/openai/baselines/issues/20

```

```

history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verb

```

```

↳

```

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 2s 39us/sample - loss: 1.6113
Epoch 2/20
60000/60000 [=====] - 1s 13us/sample - loss: 0.9672
Epoch 3/20
60000/60000 [=====] - 1s 13us/sample - loss: 0.7665
Epoch 4/20
60000/60000 [=====] - 1s 13us/sample - loss: 0.6677
Epoch 5/20
60000/60000 [=====] - 1s 13us/sample - loss: 0.6077
Epoch 6/20
60000/60000 [=====] - 1s 14us/sample - loss: 0.5667
Epoch 7/20
60000/60000 [=====] - 1s 14us/sample - loss: 0.5367
Epoch 8/20
60000/60000 [=====] - 1s 13us/sample - loss: 0.5137
Epoch 9/20
60000/60000 [=====] - 1s 13us/sample - loss: 0.4952
Epoch 10/20
60000/60000 [=====] - 1s 13us/sample - loss: 0.4800
Epoch 11/20
60000/60000 [=====] - 1s 14us/sample - loss: 0.4672
Epoch 12/20
60000/60000 [=====] - 1s 14us/sample - loss: 0.4562
Epoch 13/20
60000/60000 [=====] - 1s 13us/sample - loss: 0.4468
Epoch 14/20
60000/60000 [=====] - 1s 13us/sample - loss: 0.4384
Epoch 15/20
60000/60000 [=====] - 1s 13us/sample - loss: 0.4310
Epoch 16/20
60000/60000 [=====] - 1s 13us/sample - loss: 0.4244
Epoch 17/20
60000/60000 [=====] - 1s 13us/sample - loss: 0.4184
Epoch 18/20
60000/60000 [=====] - 1s 13us/sample - loss: 0.4130
Epoch 19/20
60000/60000 [=====] - 1s 13us/sample - loss: 0.4080
Epoch 20/20
60000/60000 [=====] - 1s 13us/sample - loss: 0.4034
```

```
score = model.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])
```

```

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch)


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

☞ Test score: 0.38070512788295746
Test accuracy: 0.9



▼ 2 Model A : 2 Hidden Laves.

▼ 2.1 MLP + ReLU activation + Adam Optimizer

```

# Multilayer perceptron

model_relu = Sequential()
model_relu.add(Dense(256, activation='relu', input_shape=(input_dim,)))
model_relu.add(Dense(128, activation='relu'))
model_relu.add(Dense(output_dim, activation='softmax'))

model_relu.summary()

```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 256)	200960
dense_2 (Dense)	(None, 128)	32896
dense_3 (Dense)	(None, 10)	1290
Total params: 235,146		
Trainable params: 235,146		
Non-trainable params: 0		

```
model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['ac  
history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,
```

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 1s 17us/sample - loss: 0.3368
Epoch 2/20
60000/60000 [=====] - 1s 15us/sample - loss: 0.1285
Epoch 3/20
60000/60000 [=====] - 1s 16us/sample - loss: 0.0840
Epoch 4/20
60000/60000 [=====] - 1s 16us/sample - loss: 0.0615
Epoch 5/20
60000/60000 [=====] - 1s 15us/sample - loss: 0.0459
Epoch 6/20
60000/60000 [=====] - 1s 15us/sample - loss: 0.0360
Epoch 7/20
60000/60000 [=====] - 1s 15us/sample - loss: 0.0257
Epoch 8/20
60000/60000 [=====] - 1s 15us/sample - loss: 0.0211
Epoch 9/20
60000/60000 [=====] - 1s 15us/sample - loss: 0.0154
Epoch 10/20
60000/60000 [=====] - 1s 16us/sample - loss: 0.0123
Epoch 11/20
60000/60000 [=====] - 1s 15us/sample - loss: 0.0095
Epoch 12/20
60000/60000 [=====] - 1s 16us/sample - loss: 0.0095
Epoch 13/20
60000/60000 [=====] - 1s 15us/sample - loss: 0.0106
Epoch 14/20
60000/60000 [=====] - 1s 16us/sample - loss: 0.0064
Epoch 15/20
60000/60000 [=====] - 1s 15us/sample - loss: 0.0054
Epoch 16/20
60000/60000 [=====] - 1s 16us/sample - loss: 0.0066
Epoch 17/20
60000/60000 [=====] - 1s 15us/sample - loss: 0.0046
Epoch 18/20
60000/60000 [=====] - 1s 15us/sample - loss: 0.0111
Epoch 19/20
60000/60000 [=====] - 1s 16us/sample - loss: 0.0120
Epoch 20/20
60000/60000 [=====] - 1s 16us/sample - loss: 0.0044
```

```
score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])
```

```
fig,ax = plt.subplots(1,1)
```

```

ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoc


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

➡ Test score: 0.08297917389905415
Test accuracy: 0.9815



```

w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)


fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')

```

```
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



▼ 2.2 MLP + ReLU activation + Adam Optimizer (Batch Normalization)

```
model_relu = Sequential()
model_relu.add(Dense(256, activation='relu', input_shape=(input_dim,), kernel_init
model_relu.add(BatchNormalization())

model_relu.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=
model_relu.add(BatchNormalization())

model_relu.add(Dense(output_dim, activation='softmax'))

model_relu.summary()
```

↗

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/

Model: "sequential_2"

Layer (type)	Output Shape	Param #
dense_4 (Dense)	(None, 256)	200960
batch_normalization (Batch Normalization)	(None, 256)	1024

pile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

l_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va

☞ Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 2s 27us/sample - loss: 0.2629
Epoch 2/20
60000/60000 [=====] - 1s 24us/sample - loss: 0.0949
Epoch 3/20
60000/60000 [=====] - 1s 24us/sample - loss: 0.0590
Epoch 4/20
60000/60000 [=====] - 1s 24us/sample - loss: 0.0401
Epoch 5/20
60000/60000 [=====] - 1s 24us/sample - loss: 0.0264
Epoch 6/20
60000/60000 [=====] - 1s 23us/sample - loss: 0.0192
Epoch 7/20
60000/60000 [=====] - 1s 23us/sample - loss: 0.0139
Epoch 8/20
60000/60000 [=====] - 1s 23us/sample - loss: 0.0111
Epoch 9/20
60000/60000 [=====] - 1s 24us/sample - loss: 0.0090
Epoch 10/20
60000/60000 [=====] - 1s 23us/sample - loss: 0.0087
Epoch 11/20
60000/60000 [=====] - 1s 23us/sample - loss: 0.0096
Epoch 12/20
60000/60000 [=====] - 1s 23us/sample - loss: 0.0094
Epoch 13/20
60000/60000 [=====] - 1s 23us/sample - loss: 0.0077
Epoch 14/20
60000/60000 [=====] - 1s 24us/sample - loss: 0.0066
Epoch 15/20
60000/60000 [=====] - 1s 23us/sample - loss: 0.0062
Epoch 16/20
60000/60000 [=====] - 1s 23us/sample - loss: 0.0062
Epoch 17/20
60000/60000 [=====] - 1s 23us/sample - loss: 0.0067
Epoch 18/20
60000/60000 [=====] - 1s 23us/sample - loss: 0.0042
Epoch 19/20
60000/60000 [=====] - 1s 24us/sample - loss: 0.0036
Epoch 20/20
60000/60000 [=====] - 1s 23us/sample - loss: 0.0034
```

```
score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])
```



```

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoc


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

☞ Test score: 0.0921328787419421
Test accuracy: 0.9782



```

w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)


fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)

```

```
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')
```

```
plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



▼ 2.3 MLP + ReLU activation + Adam Optimizer (Dropout)

```
from tensorflow.keras.layers import Dense, Dropout
model_relu = Sequential()
model_relu.add(Dense(256, activation='relu', input_shape=(input_dim,)))
model_relu.add(Dropout(0.5))
model_relu.add(Dense(128, activation='relu'))
model_relu.add(Dropout(0.5))
model_relu.add(Dense(output_dim, activation='softmax'))
```

```
model_relu.summary()
```



Model: "sequential_3"

Layer (type)	Output Shape	Param #
--------------	--------------	---------

```
model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['ac
```

```
history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,
```

```
☐ validate on 10000 samples
```

```

=====] - 1s 16us/sample - loss: 0.6157 - acc: 0.8078 - val_loss:
=====] - 1s 14us/sample - loss: 0.2708 - acc: 0.9227 - val_loss:
=====] - 1s 14us/sample - loss: 0.2128 - acc: 0.9383 - val_loss:
=====] - 1s 14us/sample - loss: 0.1786 - acc: 0.9478 - val_loss:
=====] - 1s 14us/sample - loss: 0.1581 - acc: 0.9539 - val_loss:
=====] - 1s 14us/sample - loss: 0.1410 - acc: 0.9586 - val_loss:
=====] - 1s 14us/sample - loss: 0.1271 - acc: 0.9635 - val_loss:
=====] - 1s 15us/sample - loss: 0.1178 - acc: 0.9646 - val_loss:
=====] - 1s 15us/sample - loss: 0.1107 - acc: 0.9669 - val_loss:
=====] - 1s 14us/sample - loss: 0.1069 - acc: 0.9681 - val_loss:
=====] - 1s 14us/sample - loss: 0.1029 - acc: 0.9698 - val_loss:
=====] - 1s 14us/sample - loss: 0.0941 - acc: 0.9719 - val_loss:
=====] - 1s 15us/sample - loss: 0.0919 - acc: 0.9721 - val_loss:
=====] - 1s 15us/sample - loss: 0.0861 - acc: 0.9743 - val_loss:
=====] - 1s 14us/sample - loss: 0.0822 - acc: 0.9754 - val_loss:
=====] - 1s 15us/sample - loss: 0.0807 - acc: 0.9754 - val_loss:
=====] - 1s 14us/sample - loss: 0.0792 - acc: 0.9753 - val_loss:
=====] - 1s 14us/sample - loss: 0.0754 - acc: 0.9769 - val_loss:
=====] - 1s 14us/sample - loss: 0.0697 - acc: 0.9780 - val_loss:
=====] - 1s 14us/sample - loss: 0.0705 - acc: 0.9783 - val_loss:

```

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

```

```
fig,ax = plt.subplots(1,1)
```

```

ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoc


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

☞ Test score: 0.06762538701674711
Test accuracy: 0.9807



```

w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)


fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')

```

```
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



▼ 2.4 MLP + ReLU activation + Adam Optimizer (Batch Normalization ,

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-layer-in-keras
from tensorflow.keras.layers import Dense, Dropout
from tensorflow.keras.layers import BatchNormalization
```

```
model_relu = Sequential()
model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_init
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))
model_relu.add(Dense(256, activation='relu', kernel_initializer=RandomNormal(mean=
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))
model_relu.add(Dense(output_dim, activation='softmax'))

model_relu.summary()
```

☞

Model: "sequential_11"

Layer (type)	Output Shape	Param #
dense_25 (Dense)	(None, 512)	401920
batch_normalization_12 (Batch Normalization)	(None, 512)	2048
dropout_8 (Dropout)	(None, 512)	0
dense_26 (Dense)	(None, 256)	131328
batch_normalization_13 (Batch Normalization)	(None, 256)	1024
dropout_9 (Dropout)	(None, 256)	0
dense_27 (Dense)	(None, 10)	2570
Total params: 538,890		
Trainable params: 537,354		
Non-trainable params: 1,536		

```
model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['ac
history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,

```



Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 2s 29us/sample - loss: 0.5073
Epoch 2/20
60000/60000 [=====] - 1s 24us/sample - loss: 0.2434
Epoch 3/20
60000/60000 [=====] - 1s 22us/sample - loss: 0.1902
Epoch 4/20
60000/60000 [=====] - 1s 22us/sample - loss: 0.1643
Epoch 5/20
60000/60000 [=====] - 1s 23us/sample - loss: 0.1431
Epoch 6/20
60000/60000 [=====] - 1s 23us/sample - loss: 0.1314
Epoch 7/20
60000/60000 [=====] - 1s 22us/sample - loss: 0.1193
Epoch 8/20
60000/60000 [=====] - 1s 23us/sample - loss: 0.1085
Epoch 9/20
60000/60000 [=====] - 1s 22us/sample - loss: 0.1028
Epoch 10/20
60000/60000 [=====] - 1s 23us/sample - loss: 0.0978
Epoch 11/20
60000/60000 [=====] - 2s 26us/sample - loss: 0.0905
Epoch 12/20
60000/60000 [=====] - 1s 22us/sample - loss: 0.0856
Epoch 13/20
60000/60000 [=====] - 1s 23us/sample - loss: 0.0793
Epoch 14/20
60000/60000 [=====] - 1s 22us/sample - loss: 0.0776
Epoch 15/20
60000/60000 [=====] - 1s 23us/sample - loss: 0.0751
Epoch 16/20
60000/60000 [=====] - 2s 25us/sample - loss: 0.0719
Epoch 17/20
60000/60000 [=====] - 1s 24us/sample - loss: 0.0700
Epoch 18/20
60000/60000 [=====] - 1s 25us/sample - loss: 0.0646
Epoch 19/20
60000/60000 [=====] - 1s 21us/sample - loss: 0.0633
Epoch 20/20
60000/60000 [=====] - 1s 22us/sample - loss: 0.0623
```

```
score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])
```

```
fig,ax = plt.subplots(1,1)
```

```

ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoc


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

☞ Test score: 0.06247625293707533
Test accuracy: 0.9828



```

w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')


plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')

```



```
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



▼ 3 Model B : 3 Hidden Layers

▼ 3.1 MLP + ReLU activation + Adam Optimizer

```
# Multilayer perceptron
```

```
model_relu = Sequential()
model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,)))
model_relu.add(Dense(256, activation='relu'))
model_relu.add(Dense(128, activation='relu'))
model_relu.add(Dense(output_dim, activation='softmax'))

model_relu.summary()
```

```
↳
```

Model: "sequential_4"

Layer (type)	Output Shape	Param #
dense_10 (Dense)	(None, 512)	101072

```
model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['ac
```

```
history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,
```

☞ validate on 10000 samples

```
=====] - 1s 20us/sample - loss: 0.2768 - acc: 0.9198 - val_loss:
=====] - 1s 18us/sample - loss: 0.0959 - acc: 0.9713 - val_loss:
=====] - 1s 17us/sample - loss: 0.0607 - acc: 0.9811 - val_loss:
=====] - 1s 18us/sample - loss: 0.0398 - acc: 0.9880 - val_loss:
=====] - 1s 17us/sample - loss: 0.0306 - acc: 0.9901 - val_loss:
=====] - 1s 18us/sample - loss: 0.0218 - acc: 0.9931 - val_loss:
=====] - 1s 18us/sample - loss: 0.0158 - acc: 0.9949 - val_loss:
=====] - 1s 18us/sample - loss: 0.0148 - acc: 0.9952 - val_loss:
=====] - 1s 18us/sample - loss: 0.0163 - acc: 0.9944 - val_loss:
=====] - 1s 18us/sample - loss: 0.0141 - acc: 0.9954 - val_loss:
=====] - 1s 18us/sample - loss: 0.0118 - acc: 0.9959 - val_loss:
=====] - 1s 18us/sample - loss: 0.0131 - acc: 0.9953 - val_loss:
=====] - 1s 18us/sample - loss: 0.0115 - acc: 0.9962 - val_loss:
=====] - 1s 18us/sample - loss: 0.0084 - acc: 0.9975 - val_loss:
=====] - 1s 18us/sample - loss: 0.0092 - acc: 0.9969 - val_loss:
=====] - 1s 18us/sample - loss: 0.0083 - acc: 0.9972 - val_loss:
=====] - 1s 18us/sample - loss: 0.0075 - acc: 0.9973 - val_loss:
=====] - 1s 18us/sample - loss: 0.0079 - acc: 0.9972 - val_loss:
=====] - 1s 18us/sample - loss: 0.0119 - acc: 0.9962 - val_loss:
=====] - 1s 18us/sample - loss: 0.0096 - acc: 0.9965 - val_loss:
```

```
score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])
```

```
fig,ax = plt.subplots(1,1)
```

```

ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoc


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

☞ Test score: 0.08133461526854412
Test accuracy: 0.9824



```

w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)


fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')

```

```
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



▼ 3.2 MLP + ReLU activation + Adam Optimizer (Batch Normalization)

```
model_relu = Sequential()
model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_init
model_relu.add(BatchNormalization())

model_relu.add(Dense(256, activation='relu', input_shape=(input_dim,), kernel_init
model_relu.add(BatchNormalization())

model_relu.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=
model_relu.add(BatchNormalization())

model_relu.add(Dense(output_dim, activation='softmax'))

model_relu.summary()
```

☞

Model: "sequential_5"

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_14 (Dense)	(None, 512)	401920
batch_normalization_2 (Batch Normalization)	(None, 512)	2048
dense_15 (Dense)	(None, 256)	131328
batch_normalization_3 (Batch Normalization)	(None, 256)	1024
dense_16 (Dense)	(None, 128)	32896
batch_normalization_4 (Batch Normalization)	(None, 128)	512
dense_17 (Dense)	(None, 10)	1290
=====	=====	=====
Total params: 571,018		
Trainable params: 569,226		
Non-trainable params: 1,792		
=====		

```
model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['ac
history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,
☐➔
```

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 2s 33us/sample - loss: 0.1911
Epoch 2/20
60000/60000 [=====] - 2s 29us/sample - loss: 0.0636
Epoch 3/20
60000/60000 [=====] - 2s 27us/sample - loss: 0.0386
Epoch 4/20
60000/60000 [=====] - 2s 28us/sample - loss: 0.0245
Epoch 5/20
60000/60000 [=====] - 2s 28us/sample - loss: 0.0170
Epoch 6/20
60000/60000 [=====] - 2s 27us/sample - loss: 0.0163
Epoch 7/20
60000/60000 [=====] - 2s 28us/sample - loss: 0.0162
Epoch 8/20
60000/60000 [=====] - 2s 27us/sample - loss: 0.0174
Epoch 9/20
60000/60000 [=====] - 2s 28us/sample - loss: 0.0101
Epoch 10/20
60000/60000 [=====] - 2s 28us/sample - loss: 0.0088
Epoch 11/20
60000/60000 [=====] - 2s 28us/sample - loss: 0.0078
Epoch 12/20
60000/60000 [=====] - 2s 29us/sample - loss: 0.0070
Epoch 13/20
60000/60000 [=====] - 2s 28us/sample - loss: 0.0087
Epoch 14/20
60000/60000 [=====] - 2s 29us/sample - loss: 0.0090
Epoch 15/20
60000/60000 [=====] - 2s 29us/sample - loss: 0.0090
Epoch 16/20
60000/60000 [=====] - 2s 29us/sample - loss: 0.0121
Epoch 17/20
60000/60000 [=====] - 2s 26us/sample - loss: 0.0086
Epoch 18/20
60000/60000 [=====] - 2s 27us/sample - loss: 0.0051
Epoch 19/20
60000/60000 [=====] - 2s 26us/sample - loss: 0.0042
Epoch 20/20
60000/60000 [=====] - 2s 29us/sample - loss: 0.0046
```

```
score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])
```

```
fig,ax = plt.subplots(1,1)
```

```

ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoc


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

☞ Test score: 0.07600605861382428
Test accuracy: 0.9801



```

w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)


fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')

```

```
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



▼ 3.3 MLP + ReLU activation + Adam Optimizer (Dropout)

```
from tensorflow.keras.layers import Dense, Dropout
model_relu = Sequential()
model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,)))
model_relu.add(Dropout(0.5))
model_relu.add(Dense(256, activation='relu', input_shape=(input_dim,)))
model_relu.add(Dropout(0.5))
model_relu.add(Dense(128, activation='relu'))
model_relu.add(Dropout(0.5))
model_relu.add(Dense(output_dim, activation='softmax'))

model_relu.summary()
```

↗

Model: "sequential_7"

Layer (type)	Output Shape	Param #
dense_22 (Dense)	(None, 512)	401920
dropout_5 (Dropout)	(None, 512)	0
dense_23 (Dense)	(None, 256)	131328

```
adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
in, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_
```

☞ validate on 10000 samples

```
=====] - 1s 25us/sample - loss: 0.6499 - acc: 0.7962 - val_loss:
=====] - 1s 19us/sample - loss: 0.2462 - acc: 0.9301 - val_loss:
=====] - 1s 19us/sample - loss: 0.1854 - acc: 0.9484 - val_loss:
=====] - 1s 19us/sample - loss: 0.1562 - acc: 0.9573 - val_loss:
=====] - 1s 19us/sample - loss: 0.1335 - acc: 0.9629 - val_loss:
=====] - 1s 19us/sample - loss: 0.1195 - acc: 0.9667 - val_loss:
=====] - 1s 20us/sample - loss: 0.1113 - acc: 0.9682 - val_loss:
=====] - 1s 20us/sample - loss: 0.1023 - acc: 0.9716 - val_loss:
=====] - 1s 20us/sample - loss: 0.0945 - acc: 0.9729 - val_loss:
=====] - 1s 20us/sample - loss: 0.0880 - acc: 0.9739 - val_loss:
=====] - 1s 19us/sample - loss: 0.0833 - acc: 0.9755 - val_loss:
=====] - 1s 19us/sample - loss: 0.0772 - acc: 0.9777 - val_loss:
=====] - 1s 19us/sample - loss: 0.0764 - acc: 0.9781 - val_loss:
=====] - 1s 19us/sample - loss: 0.0708 - acc: 0.9792 - val_loss:
=====] - 1s 19us/sample - loss: 0.0686 - acc: 0.9798 - val_loss:
=====] - 1s 19us/sample - loss: 0.0648 - acc: 0.9811 - val_loss:
=====] - 1s 19us/sample - loss: 0.0609 - acc: 0.9823 - val_loss:
=====] - 1s 19us/sample - loss: 0.0591 - acc: 0.9825 - val_loss:
=====] - 1s 19us/sample - loss: 0.0594 - acc: 0.9825 - val_loss:
=====] - 1s 19us/sample - loss: 0.0577 - acc: 0.9833 - val_loss:
```

```
score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
```

```

print('test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch)


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

☞ Test score: 0.06959878407377218
Test accuracy: 0.9819

```



```

w_after = model_relu.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

```

```
plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')
```

```
plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w, color='y')
plt.xlabel('Output Layer ')
plt.show()
```



▼ 3.4 MLP + ReLU activation + Adam Optimizer (Batch Normalization ,

<https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-layer-in-keras>

```
from tensorflow.keras.layers import Dense, Dropout
from tensorflow.keras.layers import BatchNormalization
```

```
model_relu = Sequential()
model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_init
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))
model_relu.add(Dense(256, activation='relu', input_shape=(input_dim,), kernel_init
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))
model_relu.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))
model_relu.add(Dense(output_dim, activation='softmax'))

model_relu.summary()
```

☞

Model: "sequential_8"

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_26 (Dense)	(None, 512)	401920
batch_normalization_5 (Batch Normalization)	(None, 512)	2048
dropout_8 (Dropout)	(None, 512)	0
dense_27 (Dense)	(None, 256)	131328
batch_normalization_6 (Batch Normalization)	(None, 256)	1024
dropout_9 (Dropout)	(None, 256)	0
dense_28 (Dense)	(None, 128)	32896
batch_normalization_7 (Batch Normalization)	(None, 128)	512
dropout_10 (Dropout)	(None, 128)	0
dense_29 (Dense)	(None, 10)	1290
=====	=====	=====
Total params: 571,018		
Trainable params: 569,226		
Non-trainable params: 1,792		

```
model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['ac
history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,
☐➔
```

validate on 10000 samples

```

=====] - 2s 39us/sample - loss: 0.6492 - acc: 0.8044 - val_loss:
=====] - 2s 30us/sample - loss: 0.2785 - acc: 0.9188 - val_loss:
=====] - 2s 29us/sample - loss: 0.2083 - acc: 0.9402 - val_loss:
=====] - 2s 31us/sample - loss: 0.1712 - acc: 0.9503 - val_loss:
=====] - 2s 30us/sample - loss: 0.1500 - acc: 0.9558 - val_loss:
=====] - 2s 31us/sample - loss: 0.1338 - acc: 0.9602 - val_loss:
=====] - 2s 31us/sample - loss: 0.1210 - acc: 0.9644 - val_loss:
=====] - 2s 32us/sample - loss: 0.1101 - acc: 0.9674 - val_loss:
=====] - 2s 29us/sample - loss: 0.0997 - acc: 0.9694 - val_loss:
=====] - 2s 31us/sample - loss: 0.0960 - acc: 0.9713 - val_loss:
=====] - 2s 30us/sample - loss: 0.0943 - acc: 0.9716 - val_loss:
=====] - 2s 29us/sample - loss: 0.0846 - acc: 0.9743 - val_loss:
=====] - 2s 29us/sample - loss: 0.0812 - acc: 0.9750 - val_loss:
=====] - 2s 31us/sample - loss: 0.0780 - acc: 0.9760 - val_loss:
=====] - 2s 30us/sample - loss: 0.0752 - acc: 0.9771 - val_loss:
=====] - 2s 30us/sample - loss: 0.0745 - acc: 0.9771 - val_loss:
=====] - 2s 29us/sample - loss: 0.0674 - acc: 0.9788 - val_loss:
=====] - 2s 28us/sample - loss: 0.0665 - acc: 0.9796 - val_loss:
=====] - 2s 30us/sample - loss: 0.0618 - acc: 0.9811 - val_loss:
=====] - 2s 29us/sample - loss: 0.0600 - acc: 0.9815 - val_loss:

```

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

```

```
fig,ax = plt.subplots(1,1)
```

```

ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoc


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

☞ Test score: 0.06073868394402962
 Test accuracy: 0.9836



```

w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)


fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')

```

```
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



▼ 4 Model C : 5 Hidden Layers

▼ 4.1 MLP + ReLU activation + Adam Optimizer

```
# Multilayer perceptron
```

```
model_relu = Sequential()
model_relu.add(Dense(1024, activation='relu', input_shape=(input_dim,)))
model_relu.add(Dense(512, activation='relu'))
model_relu.add(Dense(256, activation='relu'))
model_relu.add(Dense(128, activation='relu'))
model_relu.add(Dense(64, activation='relu'))
model_relu.add(Dense(output_dim, activation='softmax'))
```

```
model_relu.summary()
```

```
↳
```

Model: "sequential_11"

Layer (type)	Output Shape	Param #
dense_41 (Dense)	(None, 1024)	803840
dense_42 (Dense)	(None, 512)	524800
dense_43 (Dense)	(None, 256)	131328
dense_44 (Dense)	(None, 128)	32896

```
model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['ac
```

```
history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,
```

☞ Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 2s 29us/sample - loss: 0.2723
Epoch 2/20
60000/60000 [=====] - 1s 23us/sample - loss: 0.0874
Epoch 3/20
60000/60000 [=====] - 1s 22us/sample - loss: 0.0553
Epoch 4/20
60000/60000 [=====] - 1s 22us/sample - loss: 0.0409
Epoch 5/20
60000/60000 [=====] - 1s 22us/sample - loss: 0.0321
Epoch 6/20
60000/60000 [=====] - 1s 22us/sample - loss: 0.0260
Epoch 7/20
60000/60000 [=====] - 1s 22us/sample - loss: 0.0239
Epoch 8/20
60000/60000 [=====] - 1s 22us/sample - loss: 0.0205
Epoch 9/20
60000/60000 [=====] - 1s 22us/sample - loss: 0.0180
Epoch 10/20
60000/60000 [=====] - 1s 22us/sample - loss: 0.0169
Epoch 11/20
60000/60000 [=====] - 1s 23us/sample - loss: 0.0155
Epoch 12/20
60000/60000 [=====] - 1s 22us/sample - loss: 0.0123
Epoch 13/20
60000/60000 [=====] - 1s 22us/sample - loss: 0.0128
Epoch 14/20
60000/60000 [=====] - 1s 22us/sample - loss: 0.0130
Epoch 15/20
60000/60000 [=====] - 1s 22us/sample - loss: 0.0121
Epoch 16/20
60000/60000 [=====] - 1s 22us/sample - loss: 0.0104
Epoch 17/20
60000/60000 [=====] - 1s 22us/sample - loss: 0.0086
Epoch 18/20
60000/60000 [=====] - 1s 22us/sample - loss: 0.0103
Epoch 19/20
60000/60000 [=====] - 1s 22us/sample - loss: 0.0091
Epoch 20/20
60000/60000 [=====] - 1s 22us/sample - loss: 0.0088
```

```
score = model_relu.evaluate(X_test, Y_test, verbose=0)
```

```
print('Test score: ', score[0])
```



```

print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoc


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

☞ Test score: 0.10951780684681689
Test accuracy: 0.9812



```

w_after = model_relu.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

```

```
plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')
```

```
plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



▼ 4.2 MLP + ReLU activation + Adam Optimizer (Batch Normalization)

```
model_relu = Sequential()
model_relu.add(Dense(1024, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0, std=0.01)))
model_relu.add(BatchNormalization())

model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0, std=0.01)))
model_relu.add(BatchNormalization())

model_relu.add(Dense(256, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0, std=0.01)))
model_relu.add(BatchNormalization())

model_relu.add(Dense(128, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0, std=0.01)))
model_relu.add(BatchNormalization())

model_relu.add(Dense(64, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0, std=0.01)))
model_relu.add(BatchNormalization())

model_relu.add(Dense(output_dim, activation='softmax'))

model_relu.summary()
```



Model: "sequential_13"

Layer (type)	Output Shape	Param #
dense_53 (Dense)	(None, 1024)	803840
batch_normalization_13 (Batch Normalization)	(None, 1024)	4096
dense_54 (Dense)	(None, 512)	524800
batch_normalization_14 (Batch Normalization)	(None, 512)	2048
dense_55 (Dense)	(None, 256)	131328
batch_normalization_15 (Batch Normalization)	(None, 256)	1024
dense_56 (Dense)	(None, 128)	32896
batch_normalization_16 (Batch Normalization)	(None, 128)	512
dense_57 (Dense)	(None, 64)	8256
batch_normalization_17 (Batch Normalization)	(None, 64)	256
dense_58 (Dense)	(None, 10)	650

Total params: 1,509,706
 Trainable params: 1,505,738
 Non-trainable params: 3,968

```
model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['ac
history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,
```



Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 3s 53us/sample - loss: 0.1942
Epoch 2/20
60000/60000 [=====] - 2s 36us/sample - loss: 0.0679
Epoch 3/20
60000/60000 [=====] - 2s 38us/sample - loss: 0.0450
Epoch 4/20
60000/60000 [=====] - 2s 40us/sample - loss: 0.0334
Epoch 5/20
60000/60000 [=====] - 2s 38us/sample - loss: 0.0340
Epoch 6/20
60000/60000 [=====] - 2s 38us/sample - loss: 0.0269
Epoch 7/20
60000/60000 [=====] - 2s 38us/sample - loss: 0.0206
Epoch 8/20
60000/60000 [=====] - 2s 40us/sample - loss: 0.0186
Epoch 9/20
60000/60000 [=====] - 2s 36us/sample - loss: 0.0166
Epoch 10/20
60000/60000 [=====] - 2s 36us/sample - loss: 0.0189
Epoch 11/20
60000/60000 [=====] - 2s 38us/sample - loss: 0.0163
Epoch 12/20
60000/60000 [=====] - 2s 38us/sample - loss: 0.0140
Epoch 13/20
60000/60000 [=====] - 2s 36us/sample - loss: 0.0112
Epoch 14/20
60000/60000 [=====] - 2s 36us/sample - loss: 0.0100
Epoch 15/20
60000/60000 [=====] - 2s 37us/sample - loss: 0.0121
Epoch 16/20
60000/60000 [=====] - 2s 39us/sample - loss: 0.0147
Epoch 17/20
60000/60000 [=====] - 2s 39us/sample - loss: 0.0139
Epoch 18/20
60000/60000 [=====] - 2s 39us/sample - loss: 0.0075
Epoch 19/20
60000/60000 [=====] - 2s 39us/sample - loss: 0.0093
Epoch 20/20
60000/60000 [=====] - 2s 40us/sample - loss: 0.0097
```

```
score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])
```

```
fig,ax = plt.subplots(1,1)
```

```

ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoc


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

☞ Test score: 0.07663915940060397
Test accuracy: 0.9817



```

w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)


fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')

```

```
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



▼ 4.3 MLP + ReLU activation + Adam Optimizer (Dropout)

```
from tensorflow.keras.layers import Dense, Dropout
model_relu = Sequential()
model_relu.add(Dense(1024, activation='relu', input_shape=(input_dim,)))
model_relu.add(Dropout(0.5))
model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,)))
model_relu.add(Dropout(0.5))
model_relu.add(Dense(256, activation='relu', input_shape=(input_dim,)))
model_relu.add(Dropout(0.5))
model_relu.add(Dense(128, activation='relu', input_shape=(input_dim,)))
model_relu.add(Dropout(0.5))
model_relu.add(Dense(64, activation='relu'))
model_relu.add(Dropout(0.5))
model_relu.add(Dense(output_dim, activation='softmax'))

model_relu.summary()
```

☞

Model: "sequential_14"

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_59 (Dense)	(None, 1024)	803840
dropout_11 (Dropout)	(None, 1024)	0
dense_60 (Dense)	(None, 512)	524800
dropout_12 (Dropout)	(None, 512)	0
dense_61 (Dense)	(None, 256)	131328
dropout_13 (Dropout)	(None, 256)	0
dense_62 (Dense)	(None, 128)	32896
dropout_14 (Dropout)	(None, 128)	0
dense_63 (Dense)	(None, 64)	8256
dropout_15 (Dropout)	(None, 64)	0
dense_64 (Dense)	(None, 10)	650
=====	=====	=====

Total params: 1,501,770

Trainable params: 1,501,770

Non-trainable params: 0

```
model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['ac
history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,
```



validate on 10000 samples

```

=====] - 2s 35us/sample - loss: 1.0954 - acc: 0.6175 - val_loss:
=====] - 1s 24us/sample - loss: 0.3498 - acc: 0.9142 - val_loss:
=====] - 1s 24us/sample - loss: 0.2637 - acc: 0.9396 - val_loss:
=====] - 1s 24us/sample - loss: 0.2130 - acc: 0.9511 - val_loss:
=====] - 1s 24us/sample - loss: 0.1910 - acc: 0.9578 - val_loss:
=====] - 1s 23us/sample - loss: 0.1701 - acc: 0.9613 - val_loss:
=====] - 1s 24us/sample - loss: 0.1511 - acc: 0.9658 - val_loss:
=====] - 1s 23us/sample - loss: 0.1421 - acc: 0.9680 - val_loss:
=====] - 1s 24us/sample - loss: 0.1342 - acc: 0.9695 - val_loss:
=====] - 1s 24us/sample - loss: 0.1263 - acc: 0.9712 - val_loss:
=====] - 1s 24us/sample - loss: 0.1143 - acc: 0.9741 - val_loss:
=====] - 1s 23us/sample - loss: 0.1100 - acc: 0.9744 - val_loss:
=====] - 1s 23us/sample - loss: 0.1024 - acc: 0.9773 - val_loss:
=====] - 1s 23us/sample - loss: 0.0975 - acc: 0.9770 - val_loss:
=====] - 1s 23us/sample - loss: 0.0977 - acc: 0.9772 - val_loss:
=====] - 1s 23us/sample - loss: 0.0984 - acc: 0.9773 - val_loss:
=====] - 1s 23us/sample - loss: 0.0910 - acc: 0.9785 - val_loss:
=====] - 1s 24us/sample - loss: 0.0841 - acc: 0.9800 - val_loss:
=====] - 1s 23us/sample - loss: 0.0829 - acc: 0.9807 - val_loss:
=====] - 1s 24us/sample - loss: 0.0793 - acc: 0.9816 - val_loss:

```

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

```

```
fig,ax = plt.subplots(1,1)
```



```

ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoc


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

☞ Test score: 0.09344992985821433
Test accuracy: 0.9832



```

w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)


fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')

```

```
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



▼ 4.4 MLP + ReLU activation + Adam Optimizer (Batch Normalization ,

<https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalizat>
 from tensorflow.keras.layers import Dense, Dropout
 from tensorflow.keras.layers import BatchNormalization

```
model_relu = Sequential()
model_relu.add(Dense(1024, activation='relu', input_shape=(input_dim,), kernel_init
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))
model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_init
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))
model_relu.add(Dense(256, activation='relu', input_shape=(input_dim,), kernel_init
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))
model_relu.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))
model_relu.add(Dense(64, activation='relu', input_shape=(input_dim,), kernel_initi
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))
model_relu.add(Dense(output_dim, activation='softmax'))

model_relu.summary()
```

Model: "sequential_15"

Layer (type)	Output Shape	Param #
dense_65 (Dense)	(None, 1024)	803840
batch_normalization_18 (Batch Normalization)	(None, 1024)	4096
dropout_16 (Dropout)	(None, 1024)	0
dense_66 (Dense)	(None, 512)	524800
batch_normalization_19 (Batch Normalization)	(None, 512)	2048
dropout_17 (Dropout)	(None, 512)	0
dense_67 (Dense)	(None, 256)	131328
batch_normalization_20 (Batch Normalization)	(None, 256)	1024
dropout_18 (Dropout)	(None, 256)	0
dense_68 (Dense)	(None, 128)	32896
batch_normalization_21 (Batch Normalization)	(None, 128)	512
dropout_19 (Dropout)	(None, 128)	0
dense_69 (Dense)	(None, 64)	8256
batch_normalization_22 (Batch Normalization)	(None, 64)	256
dropout_20 (Dropout)	(None, 64)	0
dense_70 (Dense)	(None, 10)	650
Total params: 1,509,706		
Trainable params: 1,505,738		
Non-trainable params: 3,968		

```
optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Model:

validate on 10000 samples

```

=====] - 3s 57us/sample - loss: 0.9945 - acc: 0.6910 - val_loss:
=====] - 2s 39us/sample - loss: 0.3369 - acc: 0.9088 - val_loss:
=====] - 2s 40us/sample - loss: 0.2403 - acc: 0.9369 - val_loss:
=====] - 2s 40us/sample - loss: 0.1955 - acc: 0.9483 - val_loss:
=====] - 2s 39us/sample - loss: 0.1726 - acc: 0.9548 - val_loss:
=====] - 2s 41us/sample - loss: 0.1533 - acc: 0.9606 - val_loss:
=====] - 2s 40us/sample - loss: 0.1369 - acc: 0.9646 - val_loss:
=====] - 2s 39us/sample - loss: 0.1267 - acc: 0.9661 - val_loss:
=====] - 2s 40us/sample - loss: 0.1183 - acc: 0.9694 - val_loss:
=====] - 2s 41us/sample - loss: 0.1125 - acc: 0.9707 - val_loss:
=====] - 2s 41us/sample - loss: 0.1001 - acc: 0.9740 - val_loss:
=====] - 2s 40us/sample - loss: 0.0952 - acc: 0.9755 - val_loss:
=====] - 2s 40us/sample - loss: 0.0937 - acc: 0.9750 - val_loss:
=====] - 2s 42us/sample - loss: 0.0884 - acc: 0.9761 - val_loss:
=====] - 2s 40us/sample - loss: 0.0856 - acc: 0.9776 - val_loss:
=====] - 2s 40us/sample - loss: 0.0828 - acc: 0.9790 - val_loss:
=====] - 2s 41us/sample - loss: 0.0779 - acc: 0.9797 - val_loss:
=====] - 3s 43us/sample - loss: 0.0773 - acc: 0.9801 - val_loss:
=====] - 2s 41us/sample - loss: 0.0740 - acc: 0.9804 - val_loss:
=====] - 2s 41us/sample - loss: 0.0699 - acc: 0.9814 - val_loss:

```

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

```

```
fig,ax = plt.subplots(1,1)
```

```

ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoc


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

☞ Test score: 0.06646238846067572
Test accuracy: 0.9837



```

w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')

```